

Py-Boost custom losses

Team 13

Gradient Boosting on Decision Trees

$$R^{(t)} \approx \sum_{i=1}^m \left[g_i h_t(\mathbf{x}_i) + \frac{1}{2} d_i h_t^2(\mathbf{x}_i) \right] + \Omega(h_t)$$

$$g_i = \partial_{\hat{y}^{(t-1)}} l \left(y_i, \hat{y}^{(t-1)} \right), \quad d_i = \partial_{\hat{y}^{(t-1)}}^2 l \left(y_i, \hat{y}^{(t-1)} \right)$$

The decision:

$$\omega_j^* = -\frac{G_j}{D_j + \lambda}$$

$$G_j = \sum_{i \in I_j} g_i,$$

$$D_j = \sum_{i \in I_j} d_i$$

Problem statement

Median regression

$$MAE = \sum |y_{pred,i} - y_i|$$

$$\nabla MAE = \text{sgn}(y_{pred} - y)$$


$$\nabla^2 MAE = \text{diag}(0, \dots, 0)$$

What is a **good alternative loss function** for median regression?



Py-Boost

- GB library, supports all common features and hyperparameters
- Computationally efficient, works only on GPU, uses Cupy and Numba
- Easy to customize sampling strategy, training control, losses and metrics




CuPy



Py-Boost


- GB library, supports all common features and hyperparameters
- Computationally efficient, works only on GPU, uses Cupy and Numba
- Easy to customize sampling strategy, training control, losses and metrics



CuPy



MAE



$$MAE = \sum |y_{pred,i} - y_i|$$

$$\nabla MAE = \text{sgn}(y_{pred} - y)$$

$$\nabla^2 MAE = \text{diag}(1, \dots, 1)$$



MSLE


$$MSLE = \sum (\log(1 + y_{pred,i}) - \log(1 + y_i))^2$$

$$\log(1 + y) \rightarrow Y$$

$$MSLE = \sum (Y_{pred,i} - Y_i)^2$$



LogCosh




$$\text{LogCosh} = \sum \log(\cosh(y_{\text{pred},i} - y_i))$$

$$\nabla \text{LogCosh} = \tanh(y_{\text{pred}} - y)$$


$$\nabla^2 \text{LogCosh} = \text{diag}(\text{sech}^2(y_{\text{pred},i} - y_i))$$

Huber




$$\text{Huber}_\delta(x) = \sum \left\{ |x| \leq \delta : \frac{1}{2} x^2, \quad |x| > \delta : \delta \left(|x| - \frac{1}{2} \delta \right) \right\}$$

Pseudo Huber



$$\text{P}Huber_{\delta}(x) = \delta^2 \left(\sqrt{1 + (x/\delta)^2} - 1 \right)$$


Losses



Hyperparameter tuning

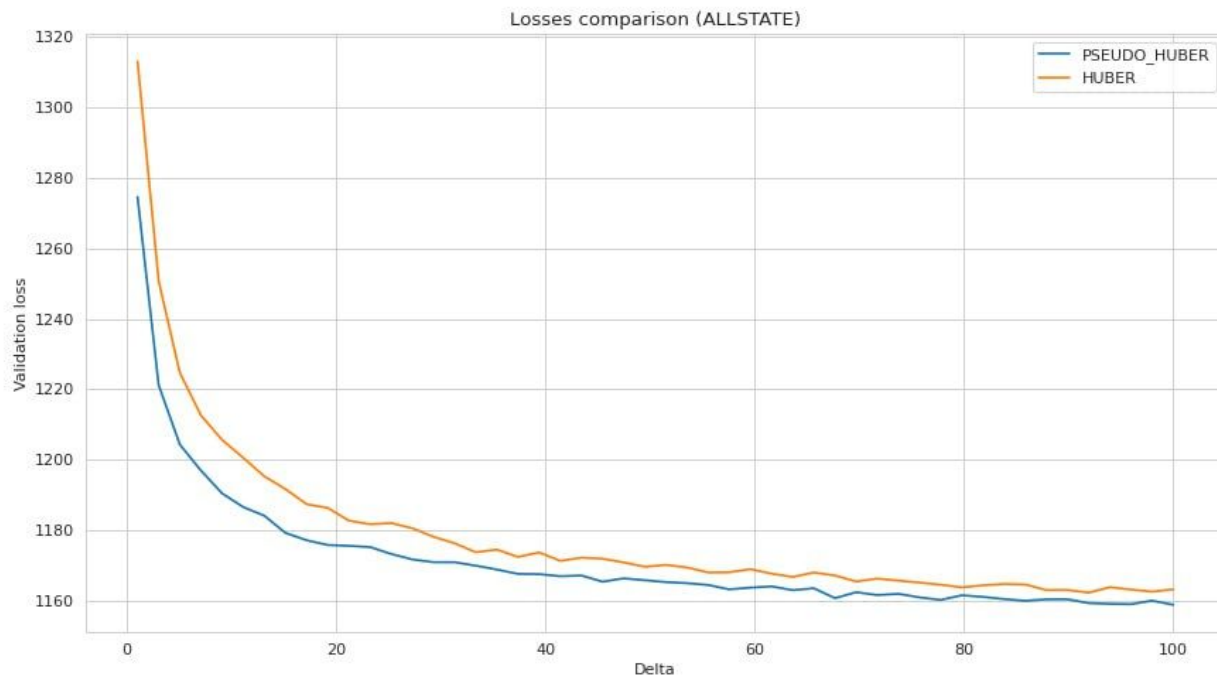
```
1 def objective(trial):
2     model = GradientBoosting(
3         loss=CustomHuberLoss(delta=trial.suggest_float("delta", 1, 20)),
4         metric=CustomMAEMetric(),
5         lr=trial.suggest_float("lr", 1e-2, 1e-1, log=True),
6         lambda_12 = trial.suggest_float("lambda_12", 1., 20.),
7         max_depth=trial.suggest_int("max_depth", 3, 10),
8         verbose = 1000
9     )
10    model.fit(X, y, eval_sets=[{'X': X_val, 'y': y_val},])
11    y_pred = model.predict(X_val)
12
13    return (np.abs(y_val - y_pred)).mean()
```

Results: Allstate dataset




	Pseudo Huber	Huber	MSLE	MAE	LogCosh
Test loss	1139	1139	1140	1915	1152

Dependence on delta: Allstate dataset




Results: House Prices dataset

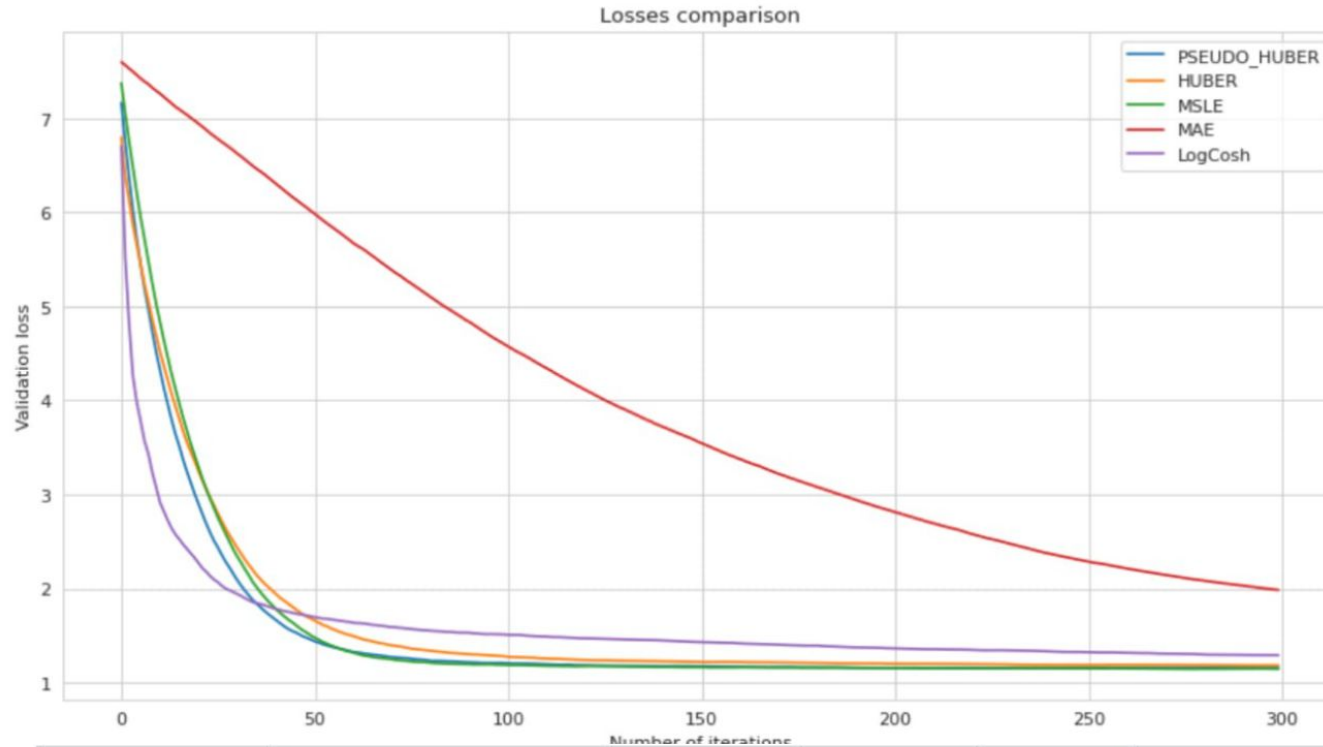


	Pseudo Huber	Huber	MSLE	MAE	LogCosh
Test loss	22818	58577	16914	62387	60336

Dependence on delta: House Prices dataset




Results: Life Expectancy dataset




	Pseudo Huber	Huber	MSLE	MAE	LogCosh
Test loss	1.17	1.16	1.13	1.20	1.34

Dependence on delta: Life Expectancy dataset



Results: Car Price Prediction dataset



	Pseudo Huber	Huber	MSLE	MAE	LogCosh
Test loss	5872	5941	5868	12691	6892

Results: Car Price Prediction dataset



Conclusion

- As expected, MAE is always a bad choice, maybe with the exception of the Life Expectancy dataset (there the scale of the target is suitable for this loss)
- MSLE always performs the best - astonishing results on the House Prices dataset illustrate this unequivocally
- The losses for Huber/Pseudo Huber mostly tend to decrease with the larger values of delta - although the plots for House Prices and Life Expectancy show some nontrivial dependencies

THANKS!



Stas Pyatkin



Dmitriy Kornilov



Danil Ivanov



Danil Gusak



Bari Khairullin